

The Stapleton Doctrine

A Personal Philosophy of AI Architecture and Programming

M. Drake Stapleton
AIEN Project

September 30, 2026
Version 1.1

English is the conversation. The verified program is the thing that was said. This paper states why the universal language of machines must be built, verified, and owned, not rented.

Abstract

On September 23, 2026, at Rails World 2026 in Austin before an audience of over a thousand developers, David Heinemeier Hansson, the creator of Ruby on Rails, retired as a professional programmer and declared: “English is a better programming language than Ruby”[1]. This paper takes that claim as stated and rejects it. English is ambiguous; two agents can read one sentence two different ways; and there is no way to prove a sentence means what you think it means. The universal thing in this author’s architecture is not a spoken language at all. It is the verified program: a small, checkable artifact, with a measured receipt, that any machine can recheck without trusting the machine that wrote it. This paper sets down the doctrine behind that architecture: how meaning is supposed to flow from suggestion to verification to evidence, how the software stack is to be stripped layer by layer down to a foundation the author owns outright, and why renting an interpreter forever is not the same as owning understanding. The argument builds on earlier marks in the same debate: Ada Lovelace on what machines can originate, Edsger Dijkstra on what tests and plain language can prove, C.A.R. Hoare on the difficulty of simplicity, and Richard Stallman on who controls whom.

1 Introduction

Watch anyone experienced work with an AI coding agent and you will see the same pattern: the human talks, the agent writes code, the human tests, the agent fixes. David Heinemeier Hansson says this pattern proves that hand-written code is over and that English is the real programming language now. His exact line, from the Rails World 2026 keynote in Austin: “English is a better programming language than Ruby”[1]. The reasoning goes like this: programming was always the translation of human ideas into machine instructions; the machine now understands human language directly; therefore the translation layer disappears, and whatever sits underneath can be a black box that nobody ever reads.

This paper agrees with one half of that argument and rejects the other.

The half that is right: talking *is* becoming the main interface. Humans should not have to hold a compiler in their heads. A person who thinks clearly but never learned to write code in detail should still be able to build serious systems by working with capable agents.

The half that is wrong: the idea that the bottom does not matter because nobody reads it. It matters more than the top. If English is the language and a rented model is the compiler

that decides what your English means, then whoever owns the model owns the meaning of your instructions. That is not the end of the meter. It is the meter moving up one layer, out of sight, where you can no longer measure it.

2 The Disagreement, On the Record

The position answered in this paper belongs to a named person and has a fixed text. On September 23, 2026, at Rails World 2026 in Austin, David Heinemeier Hansson closed a quarter century of writing code by hand and opened what he calls his retirement as a professional programmer. He has not written code by hand since about March 2026; by August he was generating around 150,000 lines a month, much of it verbose Rust he states outright he does not want to look at. In that context he said the sentence this paper answers: “English is a better programming language than Ruby” [1]. He evaluates the agent-written code “as a black box ... as any business owner in history” [3], and he names the future role the “professional maker of things” [1].

The doctrine grants the generous half of this position. Talking is becoming the interface, and Hansson is right that a careful person can now make real things through agents. It rejects the rest. A black box you own and may open is one thing; a black box whose contents live below the hand of a machine you rent, and whose meaning is settled by interpretation, is another. The difference between the two is the whole argument of this paper.

There is a precedent for answering claims by name. Alan Turing, in “Computing Machinery and Intelligence” (1950), takes the objection he attributes to Ada Lovelace, labels it “Lady Lovelace’s Objection,” and answers it by name. This paper follows that courtesy: the claim gets its speaker’s name, its date, and its exact text, and is then met on its merits [4].

The claim also has a history older than the current argument, and the history is instructive. English as a programming language is not a new idea. COBOL, designed in 1959 under the CODASYL committee at Grace Hopper’s impetus, was made to read like English on purpose, so that business managers could follow their own programs without thinking in the machine’s terms [11]. It ran, at mainframe scale, the experiment the industry now proposes to re-run: prose itself as the language of machines. The verdict was handed down within a generation, and it has not been reversed. Edsger Dijkstra, the man who spent his career arguing that correctness must be built in rather than wished in, gave the experiment its sentence: “The use of COBOL cripples the mind” [8]. In the same body of work he warns that “program testing can be used to show the presence of bugs, but never to show their absence” [6], and supplies the line that belongs on any wall of this debate: “The purpose of abstracting is not to be vague, but to create a new semantic level in which one can be absolutely precise” [7]. English is precisely vague, endlessly interpretable, and semantic levels built on it inherit that property. Hopper’s experiment served its cause and proved the point anyway: legibility is not rigor. The doctrine keeps the accessibility and puts the rigor where it belongs, in the verified artifact.

3 The Claim

English is the conversation. The program is what was actually said.

The doctrine’s central claim is simple. The candidate for a universal language of machines is not English, and it is not any language a human already speaks. It is the verified program. A program counts as a program here only if it passes every check thrown at it and carries the receipt to prove it.

“Universal” does not mean everyone can read it. It means no one has to take anyone’s word for it. A sentence can be argued with. A receipt cannot. If the artifact is checkable, any machine on earth can verify it independently. That property, and not human readability, is what makes an artifact universal.

The smallest such artifact in this architecture is called a **crumb**. A crumb is a fragment of program: drafted by an agent, checked by an independent verifier, measured for what it gains and what it costs, and filed with evidence. Programs are built the way cities are built, from the curb upward; each block laid is proven before the next one rests on it. This paper is the philosophy; the engineering rules for crumbs live in the sovereign-lineage doctrine that governs this author’s repositories.

Note what this turns down. It turns down “the code doesn’t matter because agents write it.” In this doctrine, the code is the only thing that matters. It is the evidence that the conversation ended in something real.

4 The Chain

Nothing in this architecture is believed on arrival. Everything passes through the same loop (Figure 1):

Weights suggest. Programs explain. Omega verifies. Physics authorizes. Evidence teaches.

A model’s weights propose a candidate. The candidate is restated as an explicit program, the crumb, which *explains* the proposal in the sense the loop demands: a procedure a machine can run and a checker can examine. Omega, the definitional layer, compiles the claim into something with crisp boundaries. Physics gets the last vote: only measured execution, watts drawn, bits observed, outcomes counted, counts as real. And evidence teaches, because the outcome of the measurement feeds back into the weights, so every loop ought to begin with a better suggestion than the last.




Figure 1: The verification loop. Weights suggest, programs explain, Omega verifies, Physics authorizes, evidence teaches. Evidence returns to the weights, so every pass through the loop can begin with a better suggestion than the last.

A word about the first two steps of the loop. Writing in 1843, Ada Lovelace held of Babbage’s Engine that “The Analytical Engine has no pretensions whatever to *originate* anything. It can do whatever we know how to order it to perform”[5]. The doctrine extends her line one step without breaking it: a large model can order its own suggestions, but the suggestions still originate nothing. They become something only when a program embodies them, the witness verifies them, and physics authorizes them. Suggesting and ruling are kept apart for the same reason. Lovelace’s doubt stands, answered now by proof.

A word about the verifying seat. In the working machine it is held by AEGIS. He verifies what the agents generate, including every crumb, and leaves everything as he found it: he changes nothing, decides no policy, and is not part of the generation system. A checker that rewrites the checked work is a second author, not a witness. The doctrine insists on a witness.

The loop is where meaning lives. English proposes at the top; something else, something verifiable, is what is actually said. The loop is also why this doctrine consumes less trust than the alternatives: each stage can be re-run independently, on your own machine, with your own eyes on the numbers. Trust is a tax. The loop is how you stop paying it.

5 Two Futures

The argument over the universal language reduces to two pictures of the future. Figure 2 places them side by side. They begin at the same point, a human talking to a machine, and they end in opposite places.

In the first future, English is declared the universal language and the machinery underneath is waved off as a black box. Nobody reads it; nobody needs to. But somebody pays for it, per token, forever, and somebody else decides what the words meant. The interpreter that matters is rented, and interpretation is the thing you can never stop buying.

In the second future, the conversation stays at the top, where humans are good, and verification descends into the machine, where proof lives. Programs are drafted, checked, measured, and filed. The model never gets to rule; it gets to suggest, and its suggestions are graded in the open. This future ends when there is nothing left to rent.




Figure 2: Two futures, side by side. Left: the human talks; a rented, opaque model interprets; the owner of the model owns the meaning; the meter never stops. Right: the human talks; programs are drafted, verified, and measured; the evidence is filed where anyone can check it.

The left picture is sold as freedom, “everyone can program now.” Hansson is sincere and consistent about the terms: he admits in the same breath that English is vaguer and less deterministic than the tool he gave up, and calls it a trade rather than a discount[2]. The doctrine reads it as the oldest trick in the landlord’s book: declare the foundation irrelevant, own the foundation, and collect rent on every word. The right picture asks for more work, but every bit of that work buys something: one more artifact you can prove, one more layer no one can take away. Which of the two pictures is freedom depends on where you stand: the doctrine stands at the bottom, with the receipt.

6 The Ladder

The doctrine is not content with verifying the crown of the stack. It wants the foundation. The machine a verified artifact runs on has to be ownable below the artifact, or verification is theater. So the stack is to be stripped layer by layer, down to a foundation no one else can rename, reprice, or revoke. Figure 3 shows the descent, and the doctrine’s rule for it: the ladder only goes one direction.




Figure 3: The descent. Rented services and hosted interpreters go first. Rust is the bridge. C is the next step down. Nothing finishes at the bottom until the foundation is the author’s own language. The ladder only goes one direction.

The choice to descend is the harder of the two designs C.A.R. Hoare named in his 1980 Turing lecture: “One way is to make it so simple that there are obviously no deficiencies, and the other way is to make it so complicated that there are no obvious deficiencies. The first method is far more difficult”[9]. Renting is the easy method with good lighting: a black box you can never open has no obvious deficiencies by construction, and its errors surface only when you can least afford them. Stripping the stack down to a simple bedrock you own and can read is the hard method, and Hoare’s line is the defense for doing it anyway. Difficulty is the feature, because simplicity is what proof can see.

The ladder is not an aspiration; it is a history. The whole project started when this author bounced off someone else’s language, took their layer out, and put his own in. That substitution is how AIEN began. What follows is the same move repeated on schedule: strip the hosted pieces, move down through the migration steps, and never stop before the bedrock.

Two clarifications matter. First, C is a step on this ladder, not the destination. It narrows the foundation to a minimal, exact base, but the doctrine does not end in C; it ends when the language at the bottom is mine all the way down. Second, the stripped layers are not denied, they are replaced only when a proven replacement stands under the load. The ladder descends on evidence, not on declarations. And a word about what counts as evidence along the way. Dijkstra’s warning from 1970 holds: “Program testing can be used to show the presence of bugs, but never to show their absence”[6]. A passing test says the bug you thought to look for was not there that day. A receipt is a measured result, filed with its method, ready to be re-run and re-checked by anyone. The ladder descends on the second kind of evidence, not the first.

7 Property, Not Rent

The machine that runs this doctrine has a name for the goal of the whole arrangement: **intelligence as property, not rent**. Figure 4 renders the machine in full. Read top to bottom, it is a chain of custody for meaning: thought is proposed, definition is made, verification is performed by an

independent witness, and what is verified is what the silicon executes. Nothing runs that did not pass every layer, and evidence from below feeds back up, because the loop must teach or it is decoration.

Richard Stallman stated the underlying rule a generation ago, about the software on your own disk: “The users (both individually and collectively) control the program and what it does for them. When users don’t control the program, the program controls the users”[10]. The doctrine applies the same rule one layer up, to the model that interprets your English. If the interpreter is someone else’s program under someone else’s control, then it is no metaphor to say the interpretation is not yours either. Hansson presented his black box as the end of a burden; Stallman’s line reframes it as the beginning of a custody.

This is why the doctrine and the rented-interpreter future cannot be combined by compromise. One says “the bottom does not matter.” The other says the bottom is the only thing you can truly own. Property means: the artifact on your machine can be re-run, re-measured, and re-verified by you, tomorrow, without a subscription and without anyone’s permission. Rent means someone else’s black box deciding what your instructions meant today, and charging you for the next round. The doctrine is a plain choice of property over rent, stated in an architecture rather than in a slogan.

8 The Doctrine in Tenets

The Stapleton Doctrine, stated whole:

1. **English is the conversation, not the language.** Talking is how humans steer. It is not what gets executed, not what gets verified, not what gets remembered.
2. **The verified program is the universal artifact.** Universal means checkable without trust, not readable by everyone. The crumb, with its receipt, is the currency of this architecture.
3. **Weights suggest; they do not rule.** Suggesting and ruling are different offices, and the doctrine keeps them held by different holders.
4. **Meaning descends the loop.** Suggestion, explicit program, clarified definition, measured execution, filed evidence. What arrives at the bottom is the truth of what was said at the top.
5. **The witness verifies and touches nothing.** AEGIS checks what agents generate and changes nothing, decides no policy, and belongs to no part of the generation system. Without a separate witness, verification is just more authorship.
6. **Physics is the final authorizer.** If it was not measured, it was not done. Documentation, claims, and architecture diagrams never outrank an execution receipt.
7. **Strip the stack to bedrock.** Hosted layers go first. Migration steps like C earn their place as steps. The foundation is not finished until it is owned end to end in the author’s own language.
8. **Nothing is deserved until it is replaced.** A layer is removed only when a proven replacement stands under the load. The ladder descends on evidence, not on declarations.
9. **Own the foundation or you rent the meaning.** Whoever owns the interpreter decides what your words mean. The doctrine ends when there is nothing left to rent.




Figure 4: The Sovereign Machine: the pipeline that proposes, defines, verifies, realizes, and executes, with evidence flowing back up to teach the mind. No tokens, no meters, no landlords.

10. **Eyes and ears first, industry second.** The loop must be independent of any vendor before it is convenient. Property is built before it is productized.

The Stapleton Doctrine is the philosophy behind AIEN, Omega, FORGE, AEGIS, and the Turing unit that measures their progress. English is allowed to be the conversation forever, but the machine will never be told that English is the language, because a language you cannot verify is a language you cannot own. What is said is a program; what is kept is the proof.

References

- [1] David Heinemeier Hansson, opening keynote at Rails World 2026, Austin, TX, September 23, 2026. Reported in “Ruby on Rails creator DHH says he’s done writing code by hand,” The Decoder, September 2026. <https://the-decoder.com/ruby-on-rails-creator-dhh-says-hes-done-writing-code-by-hand/>
- [2] “Ruby on Rails creator has retired from writing code by hand,” Sloptv News, September 2026. <https://sloptvnews.com/dhh-rails-world-2026-keynote-hot-takes/>
- [3] “Rails World 2026 Opening Keynote: Media Summary,” keynote notes on GitHub. <https://gist.github.com/cristoslc/5f11ff7d49a3c08a7042b8a8354a3b36>
- [4] Alan M. Turing, “Computing Machinery and Intelligence,” *Mind* **59**(236), October 1950, 433-460.
- [5] Ada Lovelace, notes on L.F. Menabrea, “Sketch of the Analytical Engine,” with notes on the memoir by the translator, *Taylor’s Scientific Memoirs* **3** (1843), 666-731. Note G carries the remark on origination.
- [6] Edsger W. Dijkstra, “Notes on Structured Programming” (EWD249), 1970. The dictum on testing appears as a corollary at the end of section 3.
- [7] Edsger W. Dijkstra, “The Humble Programmer” (EWD340), 1972 ACM Turing Award lecture, *Communications of the ACM* **15**(10), October 1972, 859-866.
- [8] Edsger W. Dijkstra, “How do we tell truths that might hurt?” (EWD498), 1975. Published in *Selected Writings on Computing: A Personal Perspective*, Springer, 1982, and in *ACM SIGPLAN Notices* **17**(5), May 1982, 13-15.
- [9] C.A.R. Hoare, “The Emperor’s Old Clothes,” 1980 ACM Turing Award lecture, *Communications of the ACM* **24**(2), February 1981, 75-83.
- [10] Richard M. Stallman, “What is Free Software?” Free Software Foundation.
- [11] COBOL, Common Business-Oriented Language. Designed 1959 by the CODASYL committee, building on Grace Hopper’s FLOW-MATIC, with the explicit goal of English-like, business-readable syntax. See e.g. <https://github.com/wallstromsimon/techshoulders/blob/HEAD/src/content/works/cobol.md>